

La Veille de Nicolas

Thursday, October 8, 2026 — evening edition

5 topic(s)

Contents

1. Claude Haiku 5.5: a small model billed as the cheapest and fastest
2. Docker Agent: build, run and share AI agents in YAML
3. Healthleap raises \$38 million for its AI that spots at-risk patients
4. New repository settings to configure pull request access
5. Claude Opus 5.5 visualizes Italo Calvino's Invisible Cities

Claude Haiku 5.5: a small model billed as the cheapest and fastest

The maker of the Claude models presents Haiku 5.5, its smallest model, designed for high-volume repetitive tasks. The post, shared on Hacker News, also details pricing, the announced results and new features for developers.

The announcement post for Claude Haiku 5.5, published on Hacker News, comes from the maker of the Claude models, who speaks of “our” models. According to the maker, Haiku 5.5 is its cheapest, fastest and most capable model to date in its category. It is aimed at high-volume, cost-sensitive repetitive tasks such as summarization, context compaction, database queries and classification. It can also serve as a sub-agent for the Opus 5.5 and Sonnet 5.5 models on coding work, and suits uses where speed matters, such as live customer support or browser use.

On pricing, the post says Haiku 5.5 costs far less than Haiku 4.5. For requests of up to 100,000 tokens, the input price falls from \$1.00 to \$0.10 per million tokens, and the output price from \$5.00 to \$0.50. Above 100,000 tokens, the reduction is 50%, according to the same document. The maker claims that on average the cost of use falls by about 75%. It states that 90% of requests sent to the previous Haiku fell into the first category, and that the new tokenizer uses slightly more tokens per task, which is taken into account in this calculation. In addition, the price of Sonnet 5.5 cache reads is halved, from \$0.20 to \$0.10 per million tokens. According to the maker, this cuts the cost of Sonnet 5.5 by about 20% on most agentic tasks.

The results cited are the maker's own. On Humanity's Last Exam without tools, it reports 45.9% for Haiku 5.5, against 10.2% for Haiku 4.5; with tools, 57.4% against 18.7%. On OSWorld 2.1, offline subset, it reports 72.4% against 15.7% for Haiku 4.5 and 48.9% for GPT-6 Luna. On Terminal-Bench 4.0, it reports 39.2%, against 70.6% for Sonnet 5.5. The post acknowledges that Sonnet 5.5 and Opus 5.5 remain better choices for complex agentic coding tasks. These figures have not been independently verified. The post points to a system card for details of the evaluation methods.

On safety, the maker states that Haiku 5.5 shows clear alignment improvements over Haiku 4.5, with fewer misaligned behaviors. Its cybersecurity safeguards are stricter than those of Haiku 4.5 but less strict than those applied to Sonnet 5.5, while still blocking penetration testing. The model is available on Amazon Web Services, Google Cloud and Microsoft Azure, under the identifier `claude-haiku-5-5`. The maker also announces a monthly API credit: \$100 a month for Max 5x subscribers, \$200 for Max 20x, and up to \$500 pooled for teams. Finally, the Python and TypeScript SDKs support computer and browser use, in beta.

Key points

- Claude Haiku 5.5 targets high-volume repetitive tasks and can serve as a sub-agent for larger models.
- According to the maker, its cost of use falls by about 75% on average compared with Haiku 4.5, figures that have not been independently verified.
- For developers: the Python and TypeScript SDKs support computer and browser use, in beta, with the identifier claude-haiku-5-5.

Sources read

Hacker News — anthropic.com — posted on HN on Oct 7, 2026 — published Oct 7, 2026

Docker Agent: build, run and share AI agents in YAML

The post published on Hacker News presents docker-agent, an extension of the Docker CLI for building and running AI agents described in a YAML file. It highlights coordination of several agents, access to MCP servers and sharing agents through an OCI registry.

The text presents docker-agent as a tool for building, running and sharing AI agents without writing code, from a declarative configuration. An agent is defined in a YAML file: it specifies the model, a description, an instruction and the list of tools. The example given uses an agent named root, based on the openai/gpt-5-mini model, with the docker:duckduckgo MCP server attached.

The multi-agent architecture makes it possible to build teams of specialized agents that split tasks among themselves automatically. Agents can use built-in tools, such as thinking, to-do tasks and memory, as well as any MCP server, whether local, remote or running in Docker. The text also mentions a document search (RAG) feature based on BM25, vector embeddings, hybrid search and result reranking.

As for providers, docker-agent does not depend on a particular model: the text names OpenAI, Anthropic, Gemini, AWS Bedrock, Mistral, xAI and Docker Model Runner, among others. Agents can be published to any OCI registry, then pulled and run elsewhere. The post also gives example commands, such as `docker agent run` to launch the default agent or `docker agent new` to generate one interactively.

For installation, three routes are given. With Docker Desktop version 4.63 or later, the extension comes preinstalled. With Homebrew, just run `brew install docker-agent`. Finally, the binaries published on GitHub Releases can be symlinked into `~/docker/cli-plugins/docker-agent`. At least one API key must also be set, or Docker Model Runner used for a local model. The text finally notes that the tool collects anonymous usage data in order to improve it.

Key points

- `docker-agent` lets you define agents in YAML, have them collaborate and share them through an OCI registry.
- It supports several model providers, including OpenAI, Anthropic and Docker Model Runner for local models.
- For developers: after installation and setting an API key, the command `docker agent run agent.yaml` launches a local configuration.

To check in the sources: three

Sources read

Hacker News — github.com — posted on HN on Oct 7, 2026 — published Sep 1, 2025 — old

Healthleap raises \$38 million for its AI that spots at-risk patients

Startup Healthleap, founded in South Africa, has raised \$38 million for an artificial intelligence that flags hospitalized patients who may need a closer look. According to TechCrunch, which had the exclusive, the funding combines a seed round and a Series A.

TechCrunch, in an article signed Ram Iyer, reports that the company Healthleap has raised \$38 million in seed and Series A funding. The \$8 million seed round was co-led by Sequoia Capital and First Round Capital. The \$30 million Series A was led by Hummingbird Ventures. The company is not disclosing its valuation.

Founded in 2022 in South Africa by the sister and brother Jemima and Josiah Meyer, the company first offered a clinical nutrition tool designed by Jemima for dietitians. It then pivoted to a more general platform, intended to spot patients admitted to hospital who may be suffering from malnutrition or acute mental confusion, two conditions often detected too late. Josiah Meyer, chief executive and co-founder, tells TechCrunch that the most telling signs are found in the notes written by caregivers.

The platform connects to the patient's electronic record and uses language models to extract mentions such as recent weight loss or difficulty swallowing. This information is then combined with lab results and vital signs in risk models. According to Josiah Meyer, the system reviews every hospitalized adult patient each night, then each morning writes a risk score into the care team's workflow, with a dashboard detailing how patients are progressing. The company states that its software makes no diagnosis and merely flags items to

examine.

According to Healthleap, the platform is deployed in more than 50 hospitals, up from 3 partners a year ago. Its customers include Penn Medicine, Cedars-Sinai, Intermountain, Houston Methodist and Emory Healthcare. Also according to the company, its revenue is said to have grown more than 10-fold over the same period, with no precise figure given. The contracts, with a 3-year term, are billed according to the number of licensed beds, with a performance-linked share. According to Josiah Meyer, every customer has seen a return on investment of at least 5 times the contract price, and in some cases more than 20 times within a year.

The research cited by TechCrunch suggests that 20% to 50% of hospitalized patients may be malnourished. Some studies link malnutrition to longer stays, impaired wound healing, infections and higher morbidity or mortality. For the University of Pennsylvania hospital, Healthleap claims that its malnutrition program generated an annualized financial impact of \$23.8 million, including \$6.3 million in additional reimbursements and \$17.5 million linked to shorter lengths of stay. These figures have not been independently verified.

The company plans to use the funds raised for engineering, product, sales and customer relations, and to add detection of new conditions. According to Josiah Meyer, programs targeting aspiration pneumonia, pressure ulcers and the risk of readmission for congestive heart failure are undergoing clinical validation. He ultimately aims for more than 40 major conditions, as well as an extension to outpatient and home care.

Key points

- Healthleap raised \$38 million, including \$8 million in seed funding and \$30 million in Series A, according to TechCrunch.
- Its tool analyzes clinical notes and structured data to flag hospitalized patients to examine, without making a diagnosis.
- According to the company, the platform is used in more than 50 hospitals, but its financial results have not been independently verified.

Sources read

TechCrunch — techcrunch.com — Oct 7, 2026

New repository settings to configure pull request access

Maintainers now have 2 settings to better control how their repositories accept contributions. A post signed Allison details how they work, their intended uses and their current limits.

The post, signed Allison, presents 2 settings that make it possible to adapt pull request handling to each project's needs. According to this text, these settings are available now for all public and private repositories, under Settings, then General, then Features.

The first setting makes it possible to disable pull requests entirely, in the same way as wikis, issues, discussions and projects. The corresponding tab then no longer appears: nobody can view existing pull requests or open new ones. The post cites mirror repositories, read-only code and projects shared publicly without contribution management as typical use cases.

The second setting limits the creation of pull requests to collaborators only, that is, users with write access. In this case the tab remains visible and anyone can read the pull requests and comment on them. Collaborators are added or removed from the Collaborators tab of the settings. The text indicates that this option helps keep contribution quality under control during critical phases of development.

The post points out two limits. The mobile app has not yet received all of the interface changes: if pull requests are disabled, the tab stays displayed there, even though nobody can create any. In addition, to temporarily restrict the activity of certain users on a public repository, temporary interaction limits remain available.

Key points

- Two settings make it possible either to disable pull requests entirely or to restrict their creation to collaborators with write access.
- For developers: these settings are found under Settings, General, Features, and collaborators are managed from the Collaborators tab.
- The mobile app still shows the pull requests tab even when they are disabled, as the interface changes are still being prepared there.

Sources read

Hacker News — github.blog — posted on HN on Oct 6, 2026 — published Feb 13, 2026 — old

Claude Opus 5.5 visualizes Italo Calvino's Invisible Cities

Piotr Migdał, who creates data visualizations and interactive explanations, gave the same prompt to GPT-6 Astra and to Claude Opus 5.5 to depict Invisible Cities. He finds Claude's result impressive, even though the time the model actually spent differs from what it claims itself.

The post, published on Hacker News, is signed Piotr Migdał. The author recounts asking two artificial intelligence models to produce an interactive visualization of Italo Calvino's collection, which presents 55 imaginary cities, each associated with an emotion or a state of mind.

The prompt asked for a visualization in three.js with pnpm, without any questions from the model, and stated that it had 6 hours of work. According to the author, GPT-6 Astra delivered a complete end-to-end result, but loaded with superfluous concepts and commentary that add visual noise. The output also reuses the style he attributes to Claude, with a beige background and numbers like 05.

Claude Opus 5.5 claimed to have used about half of the 6 hours. The author disputes this claim: according to him, the model only worked for 1 hour 25 minutes. He argues, however, that 6 sub-agents running in parallel would represent about 7 hours of agent work, which would explain the gap. The sources therefore contradict each other on the actual duration of the run.

The author specifies that his assessments are subjective. He believes GPT-6 Astra is slightly better than Fable 5.1 for the overall view, adherence to the intent of the prompt and verification that it works

correctly. Claude's result made him react enthusiastically, but no independent measurement is provided to support this judgment.

To conclude, the author invites readers to try the same prompt with other models or other tools. He thinks AI will be widely used in design, while wondering what place he himself will have in creating interactive media.

Key points

- Piotr Migdał gave the same prompt to GPT-6 Astra and Claude Opus 5.5 to visualize Calvino's *Invisible Cities*.
- Claude Opus 5.5 said it used about half of the 6 hours, whereas the author measures 1 hour 25 minutes of actual work.
- The author's assessments remain subjective and rest on no independent measurement.

Sources read

Hacker News — quesma.com — posted on HN on Oct 8, 2026 — published Oct 7, 2026